

From Markov Switching to Change Points, and Back

Ethan Goolish

1. Markov Chain Time Series

Let X be a Markov Chain with k states: S_1, \dots, S_k . It is known that we can associate with X a $(k \times k)$ transition matrix P such that given $X_t = S_i$ (meaning chain X is in state i at time t), we have that P_{ij} denotes the probability that $X_{t+1} = S_j$. However, if at time t , the state of X_t is not directly observable, but rather X_t emits some signal y_t that probabilistically depends upon the (unobservable) state of X_t , we call X a Hidden Markov Model. This is the case that we will concern ourselves with first. In specific, we associate with each state S_i a mean μ_i along with a fixed variance σ^2 so that if our chain is in state S_i at time t , we can model our emitted signal as $y_t = \mu_i + \epsilon$, where $\epsilon \stackrel{i.i.d.}{\sim} N(0, \sigma^2)$. With this, we have our Markov Chain Time Series model: Begin in a state S_1 . Using our above signal model, we obtain signal $y_1 = \mu_1 + \epsilon$. Then, we transition to a new state S_j with probability P_{1j} . We obtain our next signal $y_2 = \mu_j + \epsilon$ before transitioning to our next state $S_{j'}$ with probability $P_{jj'}$, and so forth. Thus given a $(k \times k)$ transition matrix P , a $(k \times 1)$ vector of state means S , and a variation σ^2 , we can create a time series $Y = (y_1, y_2, \dots, y_n)$ of length n . Our question is this: given a time series created in the above manner but with unknown parameters P and S , can we recover our transition probabilities and means?


2. Process

To do so, we suggest the follow procedure. In the basic case, we will assume the number of states k is known. Recall all we are given is a time series $Y = (y_1, \dots, y_n)$ of length n where each y_i has associated with it some unknown state S_j , and is drawn probabilistically as $y_i = \mu_j + \epsilon$. For example, for if we took $k = 3$ and the following parameters:

$$S = \begin{bmatrix} 5 & 0 & -5 \end{bmatrix} \quad \sigma^2 = 1 \quad P = \begin{pmatrix} 0.95 & 0.02 & 0.03 \\ 0.01 & 0.95 & 0.04 \\ 0.03 & 0.03 & 0.94 \end{pmatrix}$$


we could have the following series:

Example Time Series




Our suggested method takes advantage of the ECP change point package (TODO: cite). Within this package, we can use the *e_divisive* call to divide a time series Y into segments based on its change points. The method takes several parameters including Y , the time series in question, k , the number of estimated change points, *alpha*, the exponentiation for the distance matrix, and *min_size*, the minimum number of points allowed between change point estimations. For our method, we want to overpredict the number of change points to ensure all change points between one state S_i and another state S_j are found, and then merge the extraneous splits that separate two segments that should belong to the same state. Currently, we estimate k by assuming the transition probabilities to be no lower than 0.75, thus assuming the average time before transition to be approximately $1/(1 - 0.75) = 4$. Then we take $n/4$, where n is the total length of Y , to be a rough approximation of the number of change points, and then add $n/100$ more in order to over-estimate k appropriately. We also take *alpha* to be 2, and take *min_size* to be the smallest allowed (2) to allowed for the finest granularity. With this set of parameters, we expect a call to *e_divisive* to return segmentation similar to the following:


```
Estimated Number of Change Points: 30.0
Minimum Size before Change: 2
Time: 1.7300300598144531
Found Change Points: 30
```



Note for the most part that while the state transitions are marked by *e_divisive*, we have some excess segmentation, such as the third mark which incorrectly divides two segments from the same state S_1 , the state corresponding to $\mu_1 = 0$. Our next goal is to remove such incorrect splitting. To do this, begin by finding the mean g_i of each guessed segment G_i (which was returned by *e_divisive*). Then, sort the list of G_i incrementally by their respective g_i . We can then run *e_divisive* for a second time on our sorted G_i to group our means into estimates for which state each segment guess was in. In other words, using *e_divisive* on a sorted list of g_i groups each G_i into a guess for the segment label \hat{S}_i . If the number of states k is known, we can use $k - 1$ as the parameter for the number of change points in this second use of *e_divisive*. To continue the above example, here we plot the sorted means g_i , and then segment using the second *e_divisive* call. The segments G_i corresponding to points in the first segment will be estimated to be in state \hat{S}_0 , the segments corresponding to points in the second segment will be estimated to be in state \hat{S}_1 , and the segments corresponding to points in the third segment will be estimated to be in state \hat{S}_2 :



After assigning each G_i to the appropriate \hat{S}_i , we then do the merging – we can iterate through our original (unsorted) list of G_i , and if some G_i and G_{i+1} have the same assigned state label \hat{S}_i , then we propose a new segment G'_i such that G'_i contains the points that used to be in G_i and G_{i+1} . After doing this, we get this new estimation of the change points, which we can see is much more accurate:



Now with our fixed segmentation, we estimate our transition probabilities $P_{i,j}$ as well

as our state means μ_i . To do so, we first calculate the average stay time for each state label \hat{S}_i . By this we mean we take each updated segment G_i and observe its length. We groupby the segment labels \hat{S}_i and take the average to find the average time of stay for a given label \hat{S}_i , call it a_i . Then note that

$$a_i = \frac{1}{1 - P_{i,i}} \iff P_{i,i} = \frac{a_i - 1}{a_i}$$

and thus we can estimate our staying probabilities $P_{i,i}$. We can also estimate our off-diagonal probabilities $P_{i,j}$: We first count the number of times we transition from a state i to every other state j where $j \neq i$, call it $c_{i,j}$. Then, we know that $P_{i,j}$ is proportional to the number of times we transition from i to j over the number of times we transition from i to any other state j' . Since we know the sum across a row $\sum_j P_{i,j} = 1$, we can say

$$P_{i,j} = \frac{c_{i,j}(1 - P_{i,i})}{\sum_j c_{i,j}}$$

Alternatively, we can take a prior α_i so that $c'_{i,j} = c_{i,j} + \alpha_i$. In this basic example we take $\alpha_i = 0$, but other natural choices include $\alpha_i = 1$, $\alpha_i = 1/2$, or $\alpha_i = 1/K$ (where K is our number of states). In total, this gives us our full probability matrix P .

Then, to estimate $S = \{\mu_1, \dots, \mu_k\}$, we simply groupby our updated segment labels \hat{S}_i and take the mean to find each $\hat{\mu}_i$. Doing the above on our example dataset above yields us:

$$\hat{S} = \begin{bmatrix} -5.04134908 & -0.03278774 & 4.91458622 \end{bmatrix} \quad \hat{P} = \begin{pmatrix} 0.95384615 & 0.01538462 & 0.03076923 \\ 0.01818182 & 0.93636364 & 0.04545455 \\ 0.01298077 & 0.02163462 & 0.96538462 \end{pmatrix}$$

yielding us an accurate recovering of our original parameters.

DEPARTMENT OF STATISTICAL SCIENCE, CORNELL UNIVERSITY, ITHACA, NEW YORK, 14853
E-mail address: `efg36@cornell.edu`